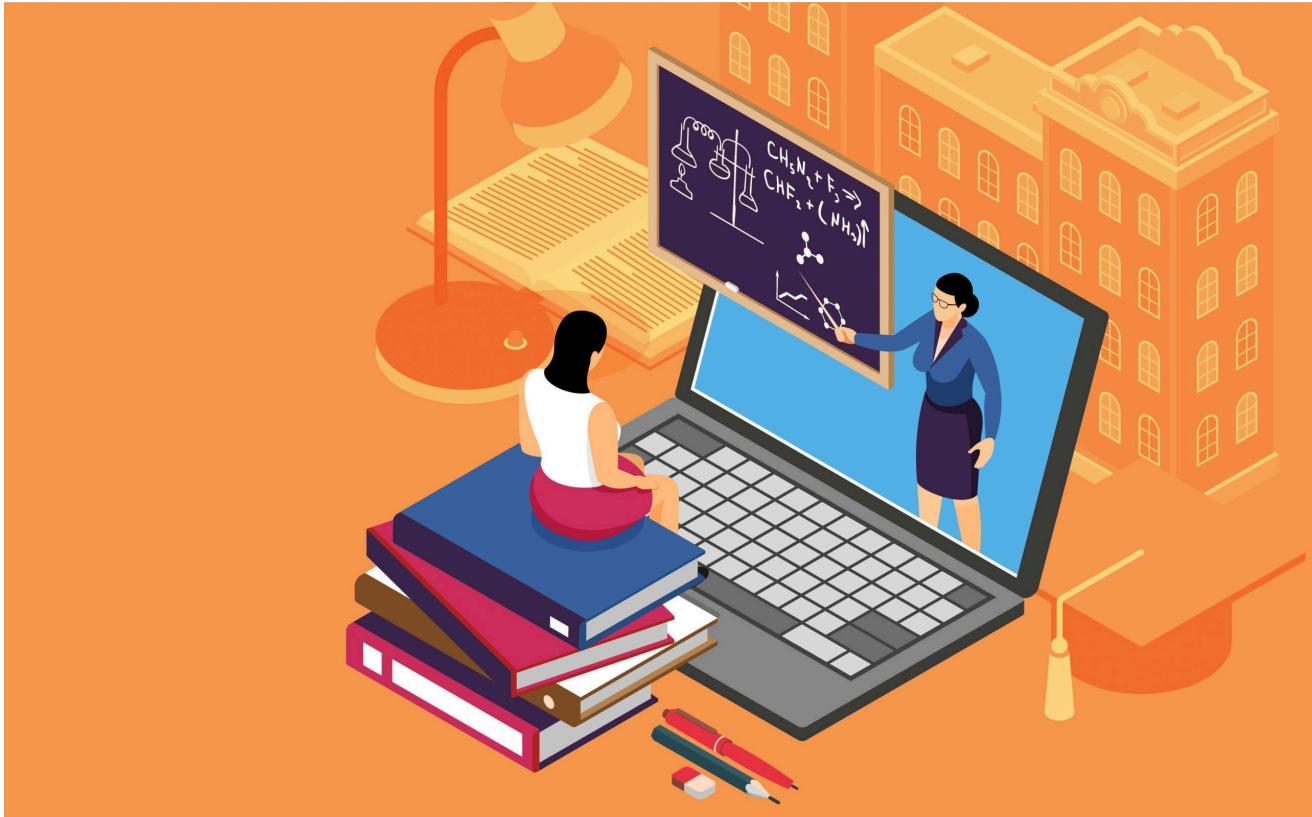


## аги Introduction



Online education services, such as Khan Academy and Coursera, provide a broader audience with access to high-quality education. On these platforms, students can learn new materials by watching a lecture, reading course material, and talking to instructors in a forum. However, one disadvantage of the online platform is that it is challenging to measure students' understanding of the course material. To deal with this issue, many online education platforms include an assessment component to ensure that students understand the core topics. The assessment component is often composed of diagnostic questions, each a multiple choice question with one correct answer. The diagnostic question is designed so that each of the incorrect answers highlights a common misconception. When students incorrectly answer the diagnostic question, it reveals the nature of their misconception and, by understanding these misconceptions, the platform can offer additional guidance to help resolve them.

у\_3 клітинки приховано

## аги 1. Download the starter code

We first install the starter code from the course website

```
!wget http://www.cs.toronto.edu/~rgrosse/courses/csc311_f20/project/final_project.zip
!unzip final_project.zip

--2023-06-19 02:28:46-- http://www.cs.toronto.edu/~rgrosse/courses/csc311_f20/project/final_project.zip
Resolving www.cs.toronto.edu (www.cs.toronto.edu)... 128.100.3.30
Connecting to www.cs.toronto.edu (www.cs.toronto.edu)|128.100.3.30|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 444833 (434K) [application/zip]
Saving to: 'final_project.zip.6'

final_project.zip.6 100%[=====] 434.41K 2.02MB/s   in 0.2s

2023-06-19 02:28:46 (2.02 MB/s) - 'final_project.zip.6' saved [444833/444833]

Archive: final_project.zip
replace starter_code/majority_vote.py? [y]es, [n]o, [A]ll, [N]one, [r]ename:
```

```
!ls starter_code/
data majority_vote.py part_a __pycache__ utils.py
```

```
!ls starter_code/data
```

```
private_test_data.csv student_meta.csv test_data.csv train_sparse.npz
question_meta.csv subject_meta.csv train_data.csv valid_data.csv
```

Let's also add the environment variable so that we can use functions provided in the starter code (`utils.py`)

```
import sys
sys.path.append('starter_code/')
```

We can now use all functions in `utils.py`!

```
from utils import *
```

## 2. Exploring the Dataset

The dataset is accessible in two ways:

- dictionary representation
- sparse matrix representation

We can use a different data format you for each learning algorithm we implement

### Dictionary Representation

Let's first load the training dataset in a dictionary

```
train_data = load_train_csv("starter_code/data")
```

The datasets loaded is a dictionary that has three keys:

- question\_id: ID of the question answered (starts from 0)
- user\_id: ID of the student who answered the question (starts from 0)
- is\_correct: Binary indicator whether the student's answer was correct (0 is incorrect, 1 is correct)

```
train_data.keys()
dict_keys(['user_id', 'question_id', 'is_correct'])
```

We can examine how many entries we have for the training data

```
len(train_data["user_id"])
```

```
56688
```

```
len(train_data["user_id"]) == len(train_data["question_id"]) == len(train_data["is_correct"])
```

```
True
```

There are in total 56688 entries in the training data. We can further see how many unique students and questions we have in the training data

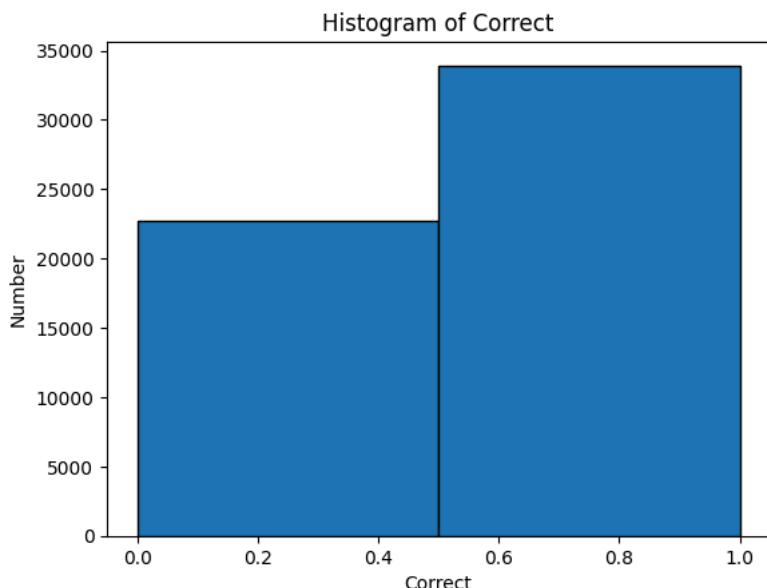
```
len(set(train_data["user_id"])), len(set(train_data["question_id"])),
(542, 1774)
```

There are in total 542 unique users and 1774 unique questions!

Next, let's observe how many correct and incorrect answers there are in the training data

```
import matplotlib.pyplot as plt

plt.hist(train_data["is_correct"], edgecolor='black', bins=2)
plt.xlabel("Correct")
plt.ylabel("Number")
plt.title("Histogram of Correct")
plt.show()
```



Correct answers ratio

```
count_one = np.count_nonzero(train_data["is_correct"])
count_all = len(train_data["is_correct"])

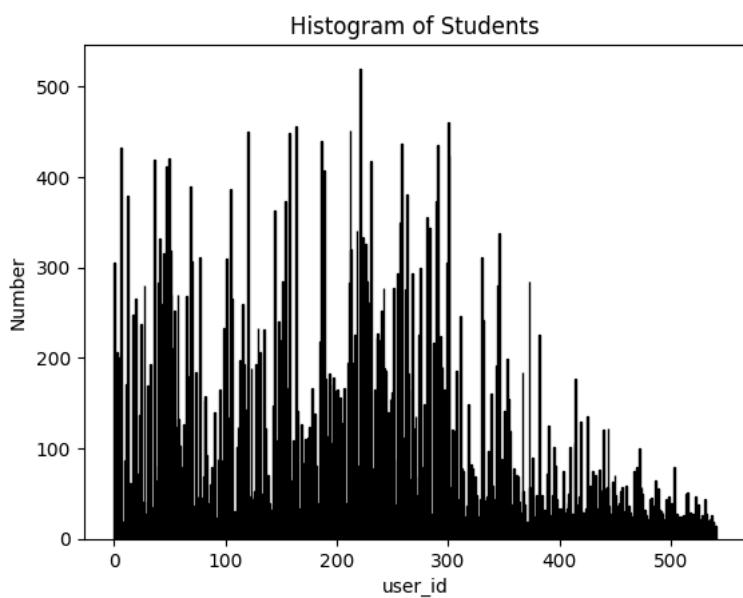
incorrect_ratio = count_one / count_all
incorrect_ratio
```

0.5984335309060118

We have more correct answers than incorrect answers in the provided dataset

We can also see how many questions each student answered

```
plt.hist(train_data["user_id"], edgecolor='black', bins=542)
plt.xlabel("user_id")
plt.ylabel("Number")
plt.title("Histogram of Students")
plt.show()
```



Some students answered more questions and some students answered less questions

```

user_answer_count = {i: 0 for i in range(0, 542)}
for u in train_data["user_id"]:
    user_answer_count[u] += 1

max(user_answer_count.values()), min(user_answer_count.values())
(520, 12)

```

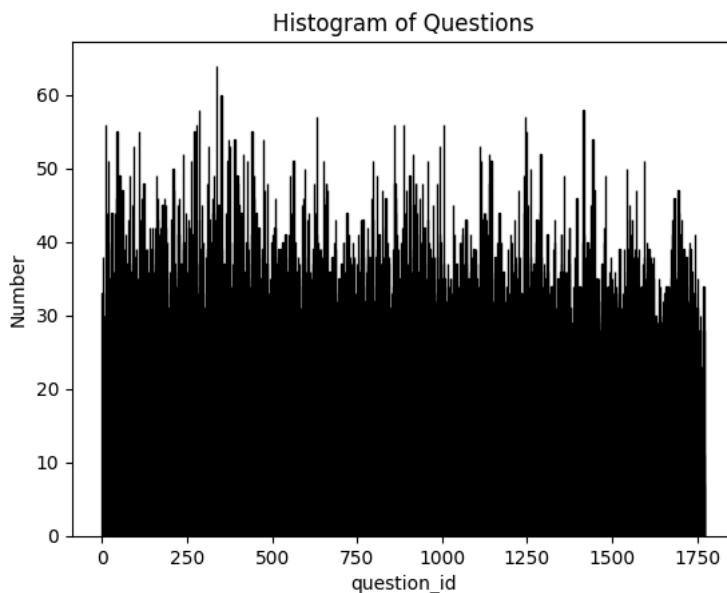
The student with most answeres answered 520 questions and the student with least answers answered 12 question

We can also examine how many times each questions was answered

```

plt.hist(train_data["question_id"], edgecolor='black', bins=1774)
plt.xlabel("question_id")
plt.ylabel("Number")
plt.title("Histogram of Questions")
plt.show()

```



```

question_answer_count = {i: 0 for i in range(0, 1774)}
for u in train_data["question_id"]:
    question_answer_count[u] += 1

max(question_answer_count.values()), min(question_answer_count.values())
(64, 5)

```

The most answered question was answered 64 times and the least answered question was answered 5 times

For now, we only looked at the training dataset. However, the way in which the validation and test dataset are stored are the same

```

valid_data = load_valid_csv("starter_code/data")
test_data = load_public_test_csv("starter_code/data")

len(valid_data["user_id"]), len(test_data["user_id"])
(7086, 3543)

```

There are 7086 and 3543 entries for each validation and test data

```

len(set(valid_data["user_id"])), len(set(valid_data["question_id"])),
(534, 1730)

len(set(test_data["user_id"])), len(set(test_data["question_id"])),
(488, 1514)

```

And there are (534, 1730) unique entries for validation data and (488, 1514) unique entries for test data

### ari Sparse matrix representation

We can also load the dataset in sparse matrix representation

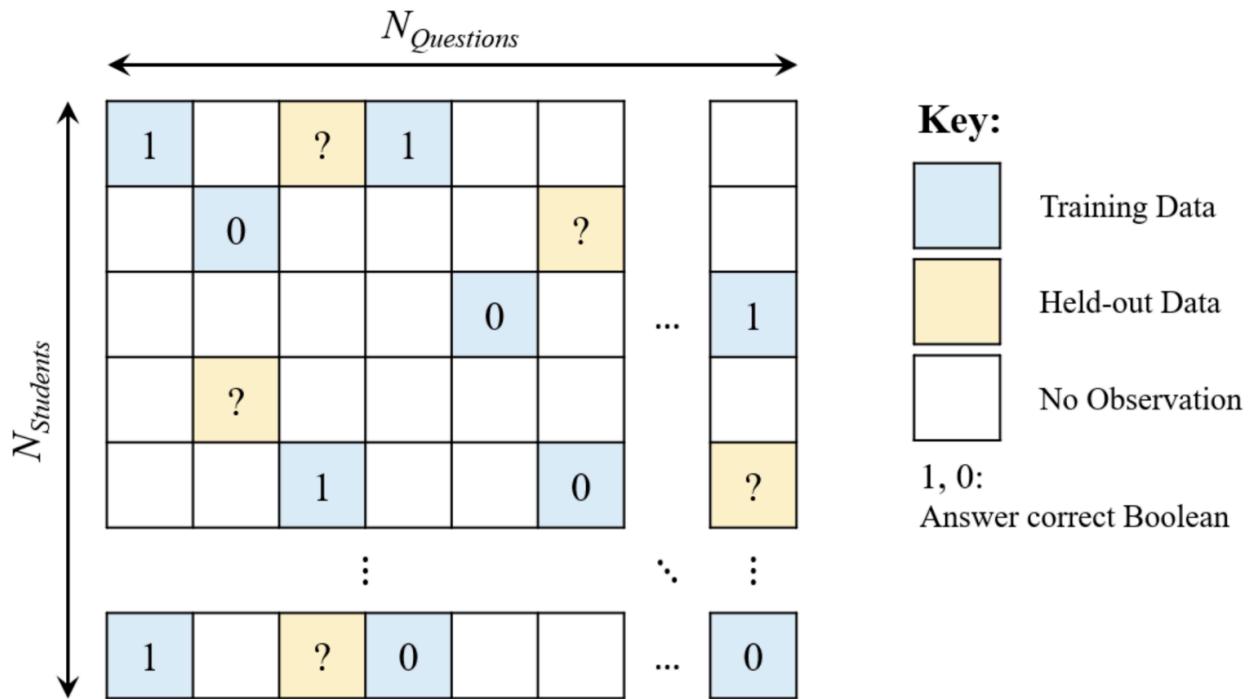


Figure 2: An example sparse matrix [1].

Each row corresponds to the user\_id and each column corresponds to the question\_id. The correct answer given a pair of (user\_id, question\_id) will have an entry 1 and an incorrect answer will have an entry 0. Answers with no observation and held-out data (that will be used for validation and test) will have an entry NaN (np.NaN)

```
matrix = load_train_sparse("starter_code/data")

matrix
<542x1774 sparse matrix of type '<class 'numpy.float64'>'  
with 938744 stored elements in Compressed Sparse Column format>

matrix.shape
(542, 1774)

matrix[0]
<1x1774 sparse matrix of type '<class 'numpy.float64'>'  
with 1620 stored elements in Compressed Sparse Column format>
```

As expected, the sparse matrix has 542 rows and 1774 columns

You may also change the type of the sparse matrix to numpy array by using `.toarray()` method

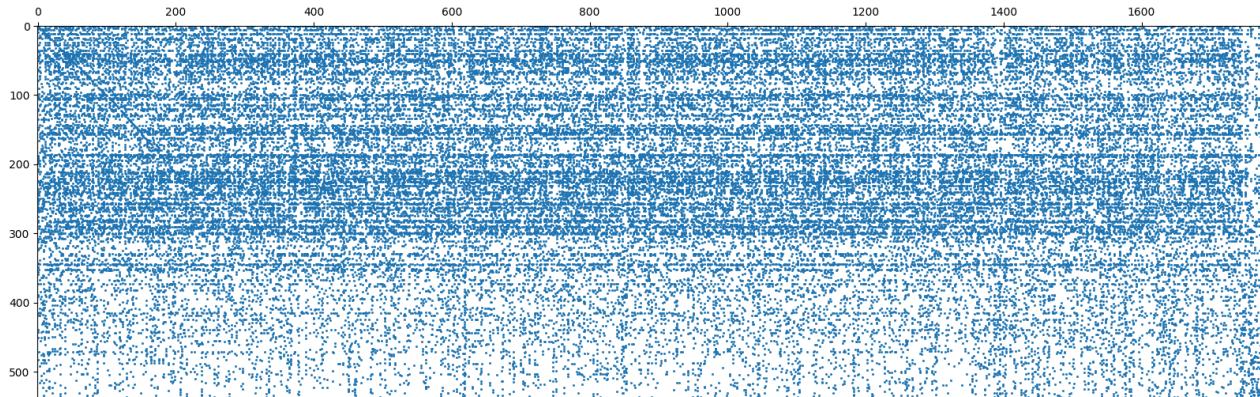
```
matrix.toarray()

array([[nan, nan, nan, ..., nan, nan, nan],
       [nan, 0., nan, ..., nan, nan, nan],
       [nan, nan, 1., ..., nan, nan, nan],
       ...,
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan]])
```

Let's visualize our sparse matrix

```
import copy
temp_matrix = copy.deepcopy(matrix).toarray()
temp_matrix[temp_matrix == 0] = 1.
temp_matrix = np.nan_to_num(temp_matrix)
plt.figure(figsize=(20, 20))
plt.spy(temp_matrix, markersize=1)
```

<matplotlib.lines.Line2D at 0x7f002f8b79a0>



We can also compute the sparsity of the matrix

```
1 - np.isnan(matrix.toarray()).sum() / np.prod(matrix.shape)
```

0.058957387770044556

This means that only 5.9% of the entries are filled in the training data

### 3. Example: prediction with majority vote

A very simple algorithm is to look whether the question was answered correctly more than 50 percent of the time. It will predict 1 if it was answered correctly more than 50 percent of the time and 0 otherwise

We first load all training, validation, and test data

```
train_data = load_train_csv("starter_code/data")
valid_data = load_valid_csv("starter_code/data")
test_data = load_public_test_csv("starter_code/data")
```

We first need to compute how many times each question was answered correctly and incorrectly

```
correct_question_map = {}
total_question_map = {}

# Count how many questions were correct.
for i, q in enumerate(train_data["question_id"]):

    if q in correct_question_map:
        if train_data["is_correct"][i] == 1:
            correct_question_map[q] += 1
    else:
        if train_data["is_correct"][i] == 1:
            correct_question_map[q] = 1
```

```

if q in total_question_map:
    total_question_map[q] += 1
else:
    total_question_map[q] = 1

```

Using these information let's predict the validation data and compute the accuracy. You may use the function `evaluate` to compute the accuracy

```

predictions = []
for i, q in enumerate(valid_data["question_id"]):
    ratio = correct_question_map[q] / float(total_question_map[q])
    # If the question was answered correctly more than half
    # of the times, predict correct.
    if ratio >= 0.5:
        predictions.append(1.)
    else:
        predictions.append(0.)
# Evaluate your model using the new prediction.
acc = evaluate(valid_data, predictions)
print("Validation Accuracy: {}".format(acc))

```

Validation Accuracy: 0.6239062941010444

We can do the same for the test data

```

predictions = []
for i, q in enumerate(test_data["question_id"]):
    ratio = correct_question_map[q] / float(total_question_map[q])
    if ratio >= 0.5:
        predictions.append(1.)
    else:
        predictions.append(0.)
acc = evaluate(test_data, predictions)
print("Test Accuracy: {}".format(acc))

```

Test Accuracy: 0.6248941574936494

We just built our first algorithm that has 62.88% accuracy on the test data! The algorithms you build should have higher accuracy than this!

## Start of personal project

By Olena Bugaiova

### My choices

Each class member selected a set of algorithms for this problem such that different students worked on different algorithms trying to achieve the best result. We compared results with accuracy

### Algorithms

I selected Matrix Factorization algorithms because the type of our data is user-item interaction data where user is a student and item is a question and the data has a lot of missing values - answers on diagnostic questions that we want to predict. We approximate missing values by reconstructing user-item interaction matrix from lower-rank matrixes of latent factors. Latent factors group students and questions in such a way that students of the same group knowing one concept covered in diagnostic questions indicate that other students of this group might know this concept too. Based on this work we can perform the task in recommending diagnostic questions

#### **Matrix Factorization**

- Singular Value Decomposition (SVD)
- Non-negative Matrix Factorization (NMF)
- Logistic Matrix Factorization (LMF)

**Bayesian Optimization** for hyperparameters tuning

### Tools

For this project I selected existing implementation from popular libraries because they are well documented and good to know

- sklearn.decomposition: SVD and NMF
- implicit: LMF
- skopt: BayesSearchCV for Bayesian Optimization
- scipy: sparse for sparse matrixes

## ari Understanding the algorithms and tools

### Singular Value Decomposition (SVD)

is the factorization of a matrix into 3 matrices

$$A = U\Sigma V^T$$

- $U$  is an  $(m \times m)$  orthogonal matrix, the left singular vectors
- $\Sigma$  is an  $(m \times n)$  nonnegative rectangular diagonal matrix, the singular values
- $V$  is an  $(n \times n)$  orthogonal matrix, the right singular vectors

The decomposition allows us to express our original matrix as a linear combination of low-rank matrices

$$\begin{matrix} & & \\ \boxed{\phantom{00}} & = & \boxed{u_1} \boxed{u_2} & \boxed{\sigma_1 \ 0} \\ A & & U & S \\ & & \boxed{0 \ \sigma_2} & \boxed{v_1^T} \\ & & & \boxed{v_2^T} & V^T \end{matrix}$$

In a practical application, we observe that only the first few, say  $K$ , singular values are large. The rest of the singular values approach zero

$$= \sigma_1 \boxed{u_1} \boxed{v_1^T} + \sigma_2 \boxed{u_2} \boxed{v_2^T}$$

### Non-negative Matrix Factorization (NMF)

A method used to factorize a **non-negative matrix**,  $X$ , into the product of two **non-negative** lower rank matrices,  $W$  and  $H$ , such that  $WH$  approximates an optimal solution of  $X$

NMF able to automatically extract sparse and **easily interpretable factors**

As an example, we take a gray-level image of a face containing  $p$  pixels, and squash the data into a single vector. Let the rows represent the pixels, and the columns each represent one image

In the case of facial images, the basis images are features such as eyes, noses, moustaches, and lips, while the columns of  $H$  indicate how much each feature is present in each image

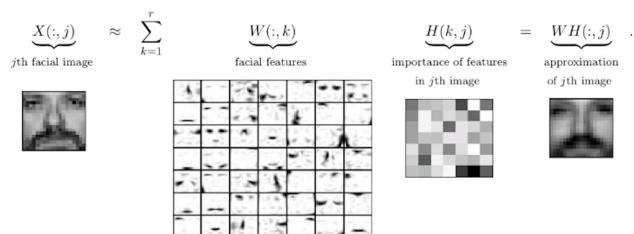


Figure 1: Decomposition of the CBCL face database, MIT Center For Biological and Computation Learning (2429 gray-level 19-by-19 pixels images) using  $r = 49$  as in [79].

The appropriate **initialization can be critical** in getting meaningful outputs

Taken from my [presentation](#) on Recommender Systems

### Logistic Matrix Factorization (LMF)

learns probabilistic distribution whether a user likes something or not. Our model factorizes the observation matrix  $R$  by 2 lower dimensional matrices  $X_{n \times f}$  and  $Y_{m \times f}$  where  $f$  is the number of latent factors... However, instead of minimizing the weighted RMSE between a binary based preference matrix and the product of  $U$  and  $V$ , we take a probabilistic approach with logistic loss.

Let  $l_{ui}$  denote the event that user  $u$  has chosen to interact with item  $i$  (user  $u$  prefers item  $i$ ). Then, we can let the probability of this event occurring be distributed according to a logistic function parameterized by the sum of the inner product of user and item latent factor vectors and user and item biases.

$$p(l_{ui} | x_u, y_i, \beta_i, \beta_j) = \frac{\exp(x_i y_i^T + \beta_u + \beta_i)}{1 + \exp(x_u y_i^T + \beta_u + \beta_i)}$$

Given this formulation, let the positive entries of our observation matrix  $r_{ui} \neq 0$  represent positive observations and the zero entries  $r_{ui} = 0$  represent negative observations

Taken from [Logistic Matrix Factorization for Implicit Feedback Data](#)

## Tools

**Sklearn.decomposition** module made for dimensionality reduction techniques mainly

**implicit** project made for recommendation algorithms for implicit feedback datasets but our data is explicit

**SciPy** is 2-D sparse array package for numeric data

## .ari Applying the algorithms and tools to our data

### Data

1. has **only 5.9%** of the entries filled with values
2. is **binary** with **zero values** for a negative event and **nan** for an absense of an event

How algorithms treat **missing values**:

- SVD is only defined for complete matrices with no missing values
- NMF in principle allows data to have missing values however NMF implementation from sklearn didn't allow nan values in the data
- LMF accepts nan values but converts them to zero
- Sparse matrixes treat nan values as zero for instance `scipy.sparse.csr_matrix.toarray()` method converts nan to zero

Matrix factorization converts data into continuous values optimizing loss functions for such values

## .ari Steps

- To add option for transformation from 0 and 1 to new binary values for instance -1, 1 in order to distinguish between negative results, posive results and missing values
- To impute missing values with fixed value like zero or with mean values for the algorithms depending on weather they can treat nan values or not and how they treat them
- To use threshold to convert predicted continues values to binary format

Also, I opened a feature request for handling nan values in NMF in sklearn

This project is focused on the classification task however with the Matrix Factorization Algorithms we could perform the diagnostic questions ranking task using metrics for ranking like Mean Average Precision at K (MAP@K), read more in [presentation](#) on Recommender Systems

## .ari 4. Import libraries

```
!pip install implicit
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: implicit in /usr/local/lib/python3.10/dist-packages (0.7.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from implicit) (1.22.4)
Requirement already satisfied: scipy>=0.16 in /usr/local/lib/python3.10/dist-packages (from implicit) (1.10.1)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from implicit) (4.65.0)
```

```
!pip install --upgrade scikit-optimize
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: scikit-optimize in /usr/local/lib/python3.10/dist-packages (0.9.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1.
Requirement already satisfied: pyyaml>=16.9 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (23.
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1
Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.10/dist-packages (from scikit-optimize) (1
Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/python3.10/dist-packages (from scikit-opti
Requirement already satisfied: PyYAML in /usr/local/lib/python3.10/dist-packages (from pyyaml>=16.9->scikit-opti
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn
```

```
pip install --upgrade scikit-learn
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.22
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.10.
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.2.
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn
```

```
import copy
import random
import pandas as pd
import numpy as np
```

```
from abc import ABC, abstractmethod
from scipy.sparse import csr_matrix
```

```
from implicit.lmf import LogisticMatrixFactorization
```

```
from sklearn.decomposition import TruncatedSVD
from sklearn.decomposition import NMF
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import balanced_accuracy_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

from sklearn.base import BaseEstimator
from sklearn.base import clone

from sklearn.utils.validation import check_is_fitted
```

```
from skopt import BayesSearchCV

from skopt.plots import plot_objective
from skopt.plots import plot_convergence
from skopt.plots import plot_evaluations
```

```
import matplotlib.pyplot as plt
%matplotlib inline
```

```
RANDOM_STATE = 22

np.random.seed(seed = RANDOM_STATE)
random.seed(RANDOM_STATE)
```

## 5. Baseline models

Trained with default hyperparameters

```
display(matrix.toarray()[:5, :4])

array([[nan, nan, nan, nan],
       [nan, 0., nan, nan],
       [nan, nan, 1., nan],
       [nan, nan, 0., 1.],
       [nan, nan, nan, nan]])
```

## Data preprocessing parameters

We can set up binary values parameters and missing value imputation strategy and then execute code on the whole section

```
DEFAULT_CORRECT_ANSWER = 1
DEFAULT_INCORRECT_ANSWER = 0
```

Set up new values for correct and incorrect answers and missing value imputation strategy

```
incorrect_answer = 0
correct_answer = 1

binary_threshold = 0.5

mean_imputation_strategy = False
fixed_imputation_strategy = True

# only for fixed imputation strategy
fixed_value_for_nan = 0.5

is_scaled = True
if correct_answer == DEFAULT_CORRECT_ANSWER and incorrect_answer == DEFAULT_INCORRECT_ANSWER:
    is_scaled = False

is_sparse = not fixed_imputation_strategy and not mean_imputation_strategy
is_negative = is_scaled and min(incorrect_answer, correct_answer) < 0
is_binary = not is_scaled and (is_sparse or (fixed_imputation_strategy and fixed_value_for_nan == 0))
```

## .ari Data transformation

```
def add_transformation(matrix, new_correct_answer = 1, new_incorrect_answer = -1):

    matrix[matrix == DEFAULT_INCORRECT_ANSWER] = 32
    matrix[matrix == DEFAULT_CORRECT_ANSWER] = new_correct_answer
    matrix[32] = new_incorrect_answer
    return matrix

def remove_transformation(matrix, scaled_correct_answer = 1, scaled_incorrect_answer = -1):

    matrix[scaled_incorrect_answer] = 32
    matrix[scaled_correct_answer] = DEFAULT_CORRECT_ANSWER
    matrix[32] = DEFAULT_INCORRECT_ANSWER

    return matrix

input_matrix = copy.deepcopy(matrix).toarray()

if is_scaled:
    input_matrix = add_transformation(input_matrix, correct_answer, incorrect_answer)
```

## .ari Missing values imputation

```
def fill_nan_with_fixed_value(matrix, fixed_value_for_nan = 0):

    matrix[np.isnan(matrix)] = fixed_value_for_nan
    return matrix

def fill_nan_with_mean(matrix):
    # obtain mean of columns
    col_mean = np.nanmean(matrix, axis = 0)

    # find indices that you need to replace
    inds = np.where(np.isnan(matrix))

    # place column means in the indices. Align the arrays using take
    matrix[inds] = np.take(col_mean, inds[1])

    return matrix, col_mean

if fixed_imputation_strategy:
    input_matrix = fill_nan_with_fixed_value(input_matrix, fixed_value_for_nan)
```

```

if mean_imputation_strategy:
    input_matrix, columns_mean = fill_nan_with_mean(input_matrix)

input_matrix[:5, :5]

array([[0.5, 0.5, 0.5, 0.5, 0.5],
       [0.5, 0., 0.5, 0.5, 0.5],
       [0.5, 0.5, 1., 0.5, 0.5],
       [0.5, 0.5, 0., 1., 0.5],
       [0.5, 0.5, 0.5, 0.5, 0.]])

```

## .arr Singular Value Decomposition

```

if is_sparse:
    matrix_svd = input_matrix
else:
    svd = TruncatedSVD(random_state = RANDOM_STATE)
    matrix_transf_svd = svd.fit_transform(input_matrix)
    matrix_svd = svd.inverse_transform(matrix_transf_svd)

print(svd.singular_values_)

[496.35572984 20.54477613]

```

## .arr Logistic Matrix Factorization

```

train_sparse_user_question = csr_matrix(
    (train_data['is_correct'], (train_data['user_id'], train_data['question_id'])))
)

lmf = LogisticMatrixFactorization(random_state = RANDOM_STATE)
lmf.fit(train_sparse_user_question)

item_factors = lmf.item_factors
user_factors = lmf.user_factors

matrix_lmf = np.dot(user_factors, item_factors.T)

0% | 0/30 [00:00<?, ?it/s]

```

lmf.item\_factors.shape

(1774, 32)

lmf.user\_factors.shape

(542, 32)

## .arr Non-Negative Matrix Factorization

```

if is_sparse or is_negative:
    matrix_nmf = input_matrix
else:
    nmf = NMF(random_state = RANDOM_STATE, verbose = 2)

    W = nmf.fit_transform(input_matrix)
    H = nmf.components_

    matrix_nmf = np.dot(W, H)

display(np.max(matrix_nmf))

Показати прихований результат

```

## .arr Converting reconstructed matrix to binary form

```

def convert_to_binary(matrix, threshold, top_value, bottom_value):
    matrix = np.where(matrix > threshold, top_value, bottom_value)
    return matrix

```

```

def convert_to_initial_form(matrix):
    top = DEFAULT_CORRECT_ANSWER
    bottom = DEFAULT_INCORRECT_ANSWER

    if is_scaled:
        top = max(correct_answer, incorrect_answer)
        bottom = min(correct_answer, incorrect_answer)

    threshold = binary_threshold
    if mean_imputation_strategy:
        threshold = columns_mean

    matrix = convert_to_binary(matrix, threshold, top, bottom)

    if is_scaled:
        matrix = remove_transformation(
            matrix,
            scaled_correct_answer = correct_answer,
            scaled_incorrect_answer = incorrect_answer
        )

    return matrix

if not is_sparse:
    matrix_svd = convert_to_initial_form(matrix_svd)

if not is_sparse and not is_negative:
    matrix_nmf = convert_to_initial_form(matrix_nmf)

matrix_lmf = convert_to_initial_form(matrix_lmf)

```

## .ari Testing predictions

Get real values

```

real_train = train_data["is_correct"]
real_valid = valid_data["is_correct"]
real_test = test_data["is_correct"]

```

Get predictions

```

def get_predictions(data, matrix):
    predictions = []

    for u, q in zip(data['user_id'], data['question_id']):
        predictions.append(matrix[u, q])

    return predictions

# predictions for SVD
predictions_train_svd = get_predictions(train_data, matrix_svd)
predictions_valid_svd = get_predictions(valid_data, matrix_svd)

# predictions for NMF
predictions_train_nmf = get_predictions(train_data, matrix_nmf)
predictions_valid_nmf = get_predictions(valid_data, matrix_nmf)

# predictions for LMF
predictions_train_lmf = get_predictions(train_data, matrix_lmf)
predictions_valid_lmf = get_predictions(valid_data, matrix_lmf)

```

Get scores

```

def get_score(real_data, predictions):
    metrics = {}

    accuracy = accuracy_score(real_data, predictions)
    balanced_accuracy = balanced_accuracy_score(real_data, predictions)
    recall = recall_score(real_data, predictions)

```

```

precision = precision_score(real_data, predictions)

f1 = f1_score(real_data, predictions)

metrics = {
    'accuracy': accuracy,
    'balanced_accuracy': balanced_accuracy,
    'recall': recall,
    'precision': precision,
    'f1_score': f1
}

return metrics

```

```

if not is_sparse:
    train_scores_svd = get_score(real_train, predictions_train_svd)
    valid_scores_svd = get_score(real_valid, predictions_valid_svd)

```

```

if not (is_sparse or is_negative):
    train_scores_nmf = get_score(real_train, predictions_train_nmf)
    valid_scores_nmf = get_score(real_valid, predictions_valid_nmf)

```

```

train_scores_lmf = get_score(real_train, predictions_train_lmf)
valid_scores_lmf = get_score(real_valid, predictions_valid_lmf)

```

```

scores = []
column_names = []
if not is_sparse:
    scores.append(train_scores_svd)
    scores.append(valid_scores_svd)
    column_names.extend(['train_svd', 'valid_svd'])

if not is_sparse and not is_negative:
    scores.append(train_scores_nmf)
    scores.append(valid_scores_nmf)
    column_names.extend(['train_nmf', 'valid_nmf'])

scores.append(train_scores_lmf)
scores.append(valid_scores_lmf)
column_names.extend(['train_lmf', 'valid_lmf'])

```

## .ari Result scores

```

scores_df = pd.DataFrame.from_dict(
    scores
)
scores_df = scores_df.T
scores_df.columns = column_names
display(scores_df.style.background_gradient(cmap = 'GnBu'))

```

	train_svd	valid_svd	train_nmf	valid_nmf	train_lmf	valid_lmf
<b>accuracy</b>	0.716042	0.692774	1.000000	0.548688	0.748395	0.570562
<b>balanced_accuracy</b>	0.695073	0.670965	1.000000	0.538180	0.705753	0.555318
<b>recall</b>	0.801586	0.779187	1.000000	0.590322	0.922356	0.630961
<b>precision</b>	0.743811	0.728371	1.000000	0.633476	0.729048	0.645984
<b>f1_score</b>	0.771619	0.752922	1.000000	0.611138	0.814388	0.638384

We can see that

- SVD model performs well even on default hyperparameters
- NMF model overfits the data because the training score reaches the maximum value while the validation score is not much better than a random guess. It happens because by default NMF uses all features as factor values
- LMF model overfits as well but not as much as NMF by taking training score as 0.74

## .ari 6. Improving models

With Bayesian Optimization

## ari Customized classes

### .ari Customized estimators

Matrix factorization estimators with interface for hyperparameters turning customized for binary data

### .ari Base Estimator

```
class BinaryMatrixFactorization(BaseEstimator, ABC):

    def __init__(
        self,
        correct_answer = 1,
        incorrect_answer = -1,
        binary_threshold = 0,
        value_for_nan = 0
    ):

        self.correct_answer = correct_answer
        self.incorrect_answer = incorrect_answer
        self.binary_threshold = binary_threshold
        self.value_for_nan = value_for_nan

        self._predicted_matrix = None

    def fit(self, data, y = None):

        matrix = self.get_input_matrix(data, self.value_for_nan)
        reconstructed_matrix = self.get_reconstructed_matrix(matrix)

        self._predicted_matrix = self.convert_to_binary(
            reconstructed_matrix,
            self.binary_threshold,
            self.correct_answer,
            self.incorrect_answer
        )

    @abstractmethod
    def get_reconstructed_matrix(self, matrix):
        pass

    def score(self, data, y = None):

        accuracy = self.get_accuracy_score(data, self._predicted_matrix)
        return accuracy

    @classmethod
    def get_input_matrix(cls, data, value_for_nan):

        """converts pandas dataframe into matrix
        Treats missing values according to the chosen strategy
        Args:
            data: given data in pandas dataframe format
            value_for_nan: fixed value to impute instead of missing values
        Returns:
            a matrix with handled missing values
        """

        if value_for_nan != 0:
            matrix = cls.get_matrix_with_nan_values(data)
            matrix = cls.fill_nan(matrix, value_for_nan)
        else:
            matrix = csr_matrix(
                (data['is_correct'], (data['user_id'], data['question_id'])))
            .toarray()

        return matrix

    @classmethod
    def get_accuracy_score(cls, data, predicted_matrix):

        predictions = cls.get_predictions(data, predicted_matrix)
        real_values = data["is_correct"]

        # we compare models based on accuracy during cross validation
        accuracy = accuracy_score(real_values, predictions)
```

```

        return accuracy

    @staticmethod
    def get_matrix_with_nan_values(data):

        """converts pandas dataframe into matrix keeping nan values
        Args:
            data: given data in pandas dataframe format
        Returns:
            a matrix with nan values
        """

        matrix_sparse = csr_matrix(
            (data['is_correct'], (data['user_id'], data['question_id'])))
        )

        # we mark all non nan values as one
        mask = csr_matrix(([1] * len(data['is_correct'])), (data['user_id'], data['question_id'])))

        # after converting sparse matrix to array, all nan values become zero
        mask = mask.toarray().astype(np.float32)

        # now we can replace zero with nan
        mask[mask == 0] = np.nan

        # get matrix of given data with nan for missing values
        matrix = matrix_sparse.toarray() * mask

        return matrix

    @staticmethod
    def fill_nan(matrix, value_for_nan):

        matrix[np.isnan(matrix)] = value_for_nan
        return matrix

    @staticmethod
    def get_predictions(data, predicted_matrix):
        predictions = []

        for u, q in zip(data['user_id'], data['question_id']):
            predictions.append(predicted_matrix[u, q])

        return predictions

    @staticmethod
    def convert_to_binary(matrix, threshold, top_value, bottom_value):

        matrix = np.where(matrix >= threshold, top_value, bottom_value)
        return matrix

```

ari SVD

```

class BinarySVD(BinaryMatrixFactorization):

    def __init__(
        self,
        n_components: int = 2,
        algorithm = 'randomized',
        n_iter: int = 5,
        n_oversamples: int = 10,
        power_iteration_normalizer = 'auto',
        random_state: int = RANDOM_STATE,
        tol: float = 0.0,
        correct_answer = 1,
        incorrect_answer = -1,
        binary_threshold = 0,
        value_for_nan = 0
    ):
        super().__init__(
            correct_answer = correct_answer,
            incorrect_answer = incorrect_answer,
            binary_threshold = binary_threshold,
            value_for_nan = value_for_nan
        )

        self.n_components = n_components
        self.algorithm = algorithm
        self.n_iter = n_iter

```

```

self.n_oversamples = n_oversamples
self.power_iteration_normalizer = power_iteration_normalizer
self.random_state = random_state
self.tol = tol

def get_reconstructed_matrix(self, matrix):
    print(type(self.n_components))

    model = TruncatedSVD(
        n_components = self.n_components,
        n_iter = self.n_iter,
        algorithm = self.algorithm,
        n_oversamples = self.n_oversamples,
        power_iteration_normalizer = self.power_iteration_normalizer,
        random_state = self.random_state,
        tol = self.tol
    )

    data_transformed = model.fit_transform(matrix)
    reconstructed_matrix = model.inverse_transform(data_transformed)

    return reconstructed_matrix

```

.ari LMF

```

class BinaryLMF(BinaryMatrixFactorization):

    def __init__(
        self,
        factors = 30,
        learning_rate = 1.0,
        regularization = 0.6,
        dtype = np.float32,
        iterations = 30,
        neg_prop = 30,
        use_gpu = False,
        num_threads = 0,
        random_state = RANDOM_STATE,
        correct_answer = 1,
        incorrect_answer = -1,
        binary_threshold = 0,
        value_for_nan = 0
    ):
        super().__init__(
            correct_answer = correct_answer,
            incorrect_answer = incorrect_answer,
            binary_threshold = binary_threshold,
            value_for_nan = value_for_nan
        )

        self.factors = factors
        self.learning_rate = learning_rate
        self.iterations = iterations
        self.regularization = regularization
        self.dtype = np.dtype(dtype)
        self.neg_prop = neg_prop
        self.random_state = random_state
        self.use_gpu = use_gpu
        self.num_threads = num_threads

    def get_reconstructed_matrix(self, matrix):

        model = LogisticMatrixFactorization(
            factors = self.factors,
            learning_rate = self.learning_rate,
            regularization = self.regularization,
            dtype = self.dtype,
            iterations = self.iterations,
            neg_prop = self.neg_prop,
            use_gpu = self.use_gpu,
            num_threads = self.num_threads,
            random_state = self.random_state
        )

        model.fit(csr_matrix(matrix))

        item_factors = model.item_factors
        user_factors = model.user_factors

```

```
reconstructed_matrix = np.dot(user_factors, item_factors.T)

return reconstructed_matrix
```

## .ari NMF

```
class BinaryNMF(BinaryMatrixFactorization):

    def __init__(
        self,
        n_components = None,
        *,
        init = None,
        solver = 'cd',
        beta_loss = 'frobenius',
        tol = 1e-4,
        max_iter = 200,
        random_state = RANDOM_STATE,
        alpha_W = 0.0,
        alpha_H = 'same',
        l1_ratio = 0.0,
        verbose = 0,
        shuffle = False,
        correct_answer = 1,
        incorrect_answer = 0,
        binary_threshold = 0.5,
        value_for_nan = 0.5
    ):
        super().__init__(
            correct_answer = correct_answer,
            incorrect_answer = incorrect_answer,
            binary_threshold = binary_threshold,
            value_for_nan = value_for_nan
        )

        self.n_components = n_components
        self.init = init
        self.beta_loss = beta_loss
        self.tol = tol
        self.max_iter = max_iter
        self.random_state = random_state
        self.alpha_W = alpha_W
        self.alpha_H = alpha_H
        self.l1_ratio = l1_ratio
        self.verbose = verbose
        self.solver = solver
        self.shuffle = shuffle

    def get_reconstructed_matrix(self, matrix):

        model = NMF(
            n_components = self.n_components,
            init = self.init,
            beta_loss = self.beta_loss,
            tol = self.tol,
            max_iter = self.max_iter,
            random_state = self.random_state,
            alpha_W = self.alpha_W,
            alpha_H = self.alpha_H,
            l1_ratio = self.l1_ratio,
            verbose = self.verbose,
            solver = self.solver,
            shuffle = self.shuffle
        )

        W = model.fit_transform(matrix);
        H = model.components_;

        reconstructed_matrix = np.dot(W, H)

        return reconstructed_matrix
```

## .ari Data Transformer

For data transformation and retrieval

```

class Singleton(object):

    def __new__(cls, *args, **kwds):
        it = cls.__dict__.get("__it__")
        if it is not None:
            return it
        cls.__it__ = it = object.__new__(cls)
        it.init(*args, **kwds)
        return it

    def init(self, *args, **kwds):
        pass

```

[Code of singleton](#) and a [post about singleton in python](#)

To create a singleton class, you subclass from Singleton; each subclass will have a single instance, no matter how many times its constructor is called. To further initialize the subclass instance, subclasses should override 'init' instead of **init** - the **init** method is called each time the constructor is called

```

class DataTransformer(Singleton):

    DEFAULT_CORRECT_ANSWER = 1
    DEFAULT_INCORRECT_ANSWER = 0

    def init(
        self
    ):
        self._binary_datasets = {}
        self._default_label = self.get_transformation_label(
            self.DEFAULT_INCORRECT_ANSWER, self.DEFAULT_CORRECT_ANSWER
        )
        self._binary_datasets[self._default_label] = self.get_default_data()

    def get_binary_data_for_cv(
        self,
        incorrect_answer = DEFAULT_INCORRECT_ANSWER,
        correct_answer = DEFAULT_CORRECT_ANSWER
    ) -> dict:
        """
        Retrieves given data converted to binary values specified in parameters
        If datasets don't exist converts default datasets to specified binary values
        Args:
            incorrect_answer: new value for incorrect answers
            correct_answer: new value for correct answers
        Returns:
            a dictionary of dataframes in format {'train':, 'valid':, 'test':, 'train_valid':}
        """
        transformation_key = self.get_transformation_label(
            incorrect_answer, correct_answer
        )

        if transformation_key not in self._binary_datasets:

            default_datasets = self._binary_datasets[self._default_label]

            requested_datasets = {}
            for datatype, dataset in default_datasets.items():
                requested_datasets[datatype] = self.transform_binary_dataset(
                    dataset, incorrect_answer, correct_answer
                )

            self._binary_datasets[transformation_key] = requested_datasets

        return self._binary_datasets[transformation_key]

    def get_default_data(self) -> dict:
        """
        Converts given train, valid and test datasets into pandas dataframe format
        Combines train and valid datasets for cross validation
        Returns:
            a dictionary of dataframes in format {'train':, 'valid':, 'test':, 'train_valid':}
        """
        datasets = {'train': train_data, 'valid': valid_data, 'test': test_data}
        dataframes = {}

        for data_type, data in datasets.items():
            df = self.convert_to_dataframe(data)

```

```

dataframes[data_type] = df

train_valid_df = pd.concat([dataframes['train'], dataframes['valid']])
train_valid_df.reset_index(drop = True, inplace = True)
dataframes['train_valid'] = train_valid_df

return dataframes

@staticmethod
def transform_binary_dataset(
    df,
    incorrect_answer,
    correct_answer,
    given_incorrect_answer = DEFAULT_INCORRECT_ANSWER,
    given_correct_answer = DEFAULT_CORRECT_ANSWER
):

    df_transformed = df.copy()
    df_transformed['is_correct'] = np.where(
        (df_transformed['is_correct'] == given_incorrect_answer),
        incorrect_answer,
        df_transformed['is_correct']
    )

    df_transformed['is_correct'] = np.where(
        (df_transformed['is_correct'] == given_correct_answer),
        correct_answer,
        df_transformed['is_correct']
    )
    return df_transformed

@staticmethod
def get_transformation_label(incorrect_answer, correct_answer) -> str:

    key = str(incorrect_answer) + '_' + str(correct_answer)
    return key

@staticmethod
def convert_to_dataframe(data):

    df = pd.DataFrame(zip(data['user_id'], data['question_id'], data['is_correct']))
    df.columns = ['user_id', 'question_id', 'is_correct']

    return df

```

## ari Hyperparameters optimizer

```

class HyperparameterOptimizer(BayesSearchCV):

    def __init__(
        self,
        model: BaseEstimator,
        hyperparameters_candidates_info: dict,
        n_iter: int = 100,
        cv: int = 3,
        return_train_score: bool = True,
        random_state: int = RANDOM_STATE,
    ):
        super().__init__(
            model,
            hyperparameters_candidates_info,
            n_iter = n_iter,
            cv = cv,
            return_train_score = return_train_score,
            random_state = random_state
        )

    def visualize_results(self):
        check_is_fitted(self, 'cv_results_')

        print("val. score: %s\n" % self.best_score_)

        print("best parameters: %s\n" % self.best_params_)
        print("best estimator: %s\n" % self.best_estimator_)

        _ = plot_convergence(self.optimizer_results_[0])
        plt.show()

        _ = plot_objective(self.optimizer_results_[0])

```

```

    plt.show()

    _ = plot_evaluations(self.optimizer_results_[0], bins = 10)
    plt.show()

    cv_results_df = pd.DataFrame(self.cv_results_)

    cv_results_df[['mean_test_score', 'mean_train_score']].plot(
        xlabel = 'number of iterations', ylabel = 'accuracy score', style=['o', 'rx']
    )

```

## .ari Model Tester

Test model with selected hyperparameters

```

class ModelTester:

    def __init__(
        self,
        model: BaseEstimator
    ):
        self._model = model

    def fit_test(self, train_data, test_data) -> float:

        self._model.fit(train_data)
        training_score = self._model.score(train_data)
        testing_score = self._model.score(test_data)

        print(f'training score: {training_score}')
        print(f'testing score: {testing_score}')

        return testing_score

```

## .ari Hyperparameters turning

What is Bayesian Optimization?

Bayesian Optimization is an optimization method that builds a probability model of the objective function of hyperparameter variables and uses it to select candidates for optimal hyperparameters to evaluate in the true objective function, [read more](#). Bayesian Optimization differs from Random Search and Grid Search in that it improves the search speed using past performances, whereas the other two methods are uniform (or independent) of past evaluations. In that sense, Bayesian Optimization is like manual search, [read more](#)

Flags for hyperparameters turning with cross validation

```

run_svd_cv = False
run_lmf_cv = True
run_nmf_cv = False

```

Hyperparameters:

- factors
- learning\_rate
- regularization
- initialization
- loss function

## .ari SVD

### .ari Binary data preprocessing

```

SVD_CORRECT_ANSWER = 1
SVD_INCORRECT_ANSWER = 0
SVD_BINARY_THRESHOLD = 0.5
SVD_MISSING_VALUE_IMPUTATION_STRATEGY = 0.5

```

```
data_for_cv = DataTransformer().get_binary_data_for_cv(
    SVD_INCORRECT_ANSWER, SVD_CORRECT_ANSWER
)
data_for_cv['train_valid'].tail()
```

	user_id	question_id	is_correct
63769	133	157	1
63770	64	197	1
63771	415	479	0
63772	115	410	0
63773	77	1273	0

.ar1 Hyperparameters tuning

```
if run_svd_cv:
    opt = HyperparameterOptimizer(
        BinarySVD(
            algorithm = 'randomized',
            n_iter = 50,
            power_iteration_normalizer = 'OR',
            random_state = RANDOM_STATE,
            correct_answer = SVD_CORRECT_ANSWER,
            incorrect_answer = SVD_INCORRECT_ANSWER,
            binary_threshold = SVD_BINARY_THRESHOLD,
            value_for_nan = SVD_MISSING_VALUE_IMPUTATION_STRATEGY
        ),
        {
            'n_components': (2, 100),
            'n_oversamples': (50, 1000)
        },
        n_iter = 100,
        cv = 3,
        return_train_score = True,
        random_state = RANDOM_STATE
    )
    opt.fit(data_for_cv['train_valid'])
```

Показати прихованій результат

```
if run_svd_cv:
    opt.visualize_results()
```

Показати прихованій результат

.ar1 Testing model with selected hyperparameters

```
if run_svd_cv:
    tester = ModelTester(opt.best_estimator_)
    tester.fit_test(data_for_cv['train'], data_for_cv['valid'])

<class 'int'>
training score: 0.7156541066892464
testing score: 0.6933389782670054
```

.ar1 LMF

.ar1 Binary data preprocessing

```
LMF_CORRECT_ANSWER = 1
LMF_INCORRECT_ANSWER = -1
LMF_BINARY_THRESHOLD = 0
LMF_MISSING_VALUE_IMPUTATION_STRATEGY = 0
```

```
data_for_cv = DataTransformer().get_binary_data_for_cv(
    LMF_INCORRECT_ANSWER, LMF_CORRECT_ANSWER
)
data_for_cv['train_valid'].tail()
```

user_id	question_id	is_correct
63769	133	157
63770	64	197
63771	415	479
63772	115	410
63773	77	1273

.аги Hyperparameters tuning

```
if run_lmf_cv:
    opt = HyperparameterOptimizer(
        BinaryLMF(
            random_state = RANDOM_STATE,
            iterations = 50,
            correct_answer = LMF_CORRECT_ANSWER,
            incorrect_answer = LMF_INCORRECT_ANSWER,
            binary_threshold = LMF_BINARY_THRESHOLD,
            value_for_nan = LMF_MISSING_VALUE_IMPUTATION_STRATEGY
        ),
        {
            'factors': (2, 400),
            'neg_prop': (1, 500),
            'learning_rate': (1e-7, 30.0, 'log-uniform'),
            'regularization': (1e-2, 150.0, 'log-uniform')
        },
        n_iter = 100,
        cv = 3,
        return_train_score = True,
        random_state = RANDOM_STATE
    )
    opt.fit(data_for_cv['train_valid'])
```

[Показати прихованій результат](#)

```
if run_lmf_cv:
    opt.visualize_results()
```

[Показати прихованій результат](#)

.аги Testing model with selected hyperparameters

```
if run_lmf_cv:
    tester = ModelTester(opt.best_estimator_)
    tester.fit_test(data_for_cv['train'], data_for_cv['valid'])

100%                                         50/50 [00:00<00:00, 155.11it/s]
training score: 0.7027766017499294
testing score: 0.6892464013547841
```

.аги NMF

.аги Binary data preprocessing

```
NMF_CORRECT_ANSWER = 1
NMF_INCORRECT_ANSWER = 0
NMF_BINARY_THRESHOLD = 0.5
NMF_MISSING_VALUE_IMPUTATION_STRATEGY = 0.5
```

```
data_for_cv = DataTransformer().get_binary_data_for_cv(
    NMF_INCORRECT_ANSWER, NMF_CORRECT_ANSWER
)
data_for_cv['train_valid'].tail()
```

	user_id	question_id	is_correct
	63769	133	157
.аги	63770	64	197
	63771	415	479
			0

```

if run_nmf_cv:
    opt = HyperparameterOptimizer(
        BinaryNMF(
            init = 'random',
            solver = 'mu',
            beta_loss = 'kullback-leibler',
            max_iter = 1000,
            alpha_H = 0,
            tol = 0.01,
            random_state = RANDOM_STATE,
            verbose = 10,
            correct_answer = NMF_CORRECT_ANSWER,
            incorrect_answer = NMF_INCORRECT_ANSWER,
            binary_threshold = NMF_BINARY_THRESHOLD,
            value_for_nan = NMF_MISSING_VALUE_IMPUTATION_STRATEGY
        ),
        {
            'n_components': (2, 300),
            'alpha_W': (1e-4, 20.0, 'log-uniform'),
            'l1_ratio': (0.0, 1.0, 'uniform')
        },
        n_iter = 280,
        cv = 3,
        return_train_score = True,
        random_state = RANDOM_STATE
    )

    # add y for stratification
    y = data_for_cv['train_valid'][["is_correct"]]
    opt.fit(data_for_cv['train_valid'], y)

```

[Показати прихованій результат](#)

```

if run_nmf_cv:
    opt.visualize_results()

```

[Показати прихованій результат](#)

.аги Testing model with selected hyperparameters

```

if run_nmf_cv:
    tester = ModelTester(opt.best_estimator_)
    tester.fit_test(data_for_cv['train'], data_for_cv['valid'])

Epoch 10 reached after 0.252 seconds, error: 187.252006
Epoch 20 reached after 0.531 seconds, error: 187.223169
training score: 0.7164126446514254
testing score: 0.6954558283940163

```

## .аги Results

Initialize models with selected hyperparameters

```

top_result_models = {}

model_lmf = BinaryLMF(
    factors = 2,
    regularization = 150,
    learning_rate = 2.25,
    neg_prop = 500,
    iterations = 50,
    random_state = RANDOM_STATE,
    correct_answer = 1,
    incorrect_answer = -1,
    binary_threshold = 0,
    value_for_nan = 0
)

top_result_models['lmf'] = {

```

```
'model': model_lmf, 'correct_answer': 1, 'incorrect_answer': -1
}
```

```
model_svd = BinarySVD(
    n_components = 2,
    n_oversamples = 530,
    algorithm = 'randomized',
    n_iter = 50,
    power_iteration_normalizer = 'OR',
    random_state = RANDOM_STATE,
    correct_answer = 1,
    incorrect_answer = 0,
    binary_threshold = 0.5,
    value_for_nan = 0.5
)

top_result_models['svd'] = {
    'model': model_svd, 'correct_answer': 1, 'incorrect_answer': 0
}
```

```
model_nmf = BinaryNMF(
    n_components = 2,
    alpha_W = 3.622,
    alpha_H = 0,
    l1_ratio = 0.048,
    init = 'random',
    solver = 'mu',
    beta_loss = 'kullback-leibler',
    max_iter = 100000,
    tol = 0.0001,
    random_state = RANDOM_STATE,
    correct_answer = 1,
    incorrect_answer = 0,
    binary_threshold = 0.5,
    value_for_nan = 0.5
)

top_result_models['nmf'] = {
    'model': model_nmf, 'correct_answer': 1, 'incorrect_answer': 0
}
```

```
test_scores = {}
for model_name, top_result_model in top_result_models.items():

    tester = ModelTester(top_result_model['model'])

    data_for_testing = DataTransformer().get_binary_data_for_cv(
        top_result_model['incorrect_answer'], top_result_model['correct_answer']
    )
    tester.fit_test(data_for_testing['train'], data_for_testing['valid'])

    test_scores[model_name] = tester.fit_test(
        data_for_testing['train_valid'], data_for_testing['test']
    )
```

```
100%                                         50/50 [00:00<00:00, 142.74it/s]
training score: 0.7027766017499294
testing score: 0.6896697713801863
100%                                         50/50 [00:00<00:00, 142.69it/s]
training score: 0.7011164424373569
testing score: 0.6841659610499576
<class 'int'>
training score: 0.7156541066892464
testing score: 0.6933389782670054
<class 'int'>
training score: 0.7134255339166431
testing score: 0.6926333615580017
training score: 0.7164126446514254
testing score: 0.6954558283940163
training score: 0.7164988866936369
testing score: 0.6954558283940163
```

```
scores = []
column_names = []

for model_name, model_score in test_scores.items():

    scores.append(model_score)
    column_names.append(model_name)
```

```
scores_df = pd.DataFrame.from_dict(
    scores
)
scores_df = scores_df.T
scores_df.columns = column_names

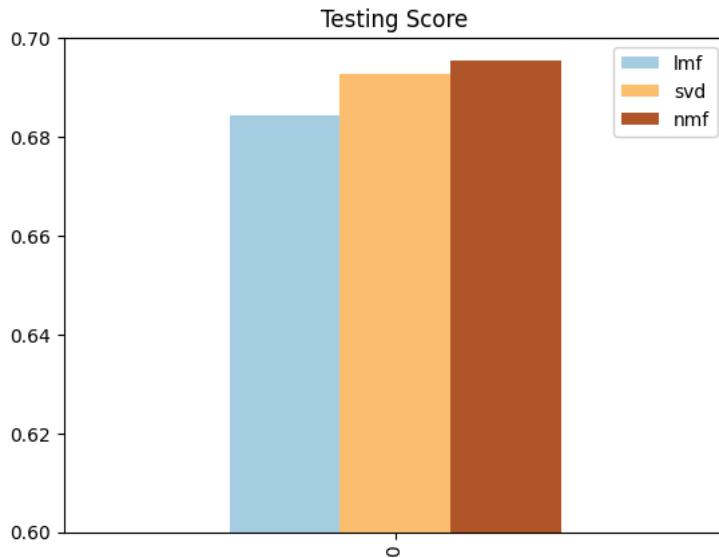
print('Testing scores:')
display(scores_df.style.background_gradient(cmap = 'GnBu'))
```

Testing scores:

	lmf	svd	nmf
0	0.684166	0.692633	0.695456

```
scores_df.plot.bar(ylim = [0.6, 0.7], cmap = 'Paired', title='Testing Score')
```

<Axes: title={'center': 'Testing Score'}>



## 7. NMF accepting missing values

Following code implemented for the "Handling NaNs in NMF" task on sklearn

Trying to solve missing values problem I found this git request that was frozen for a period of time

[Add support to missing values in NMF](#)

Then I looked into code and found parameters responsible for turning on handling missing values. But It wasn't clear if there is any implementation behind. I decided to open a new feature request which was handled soon and is awaiting approval for a merge

[Handling NaNs in NMF](#)

I downloaded source code from the feature request [ENH add support to missing values in NMF](#) from nmf\_missing branch

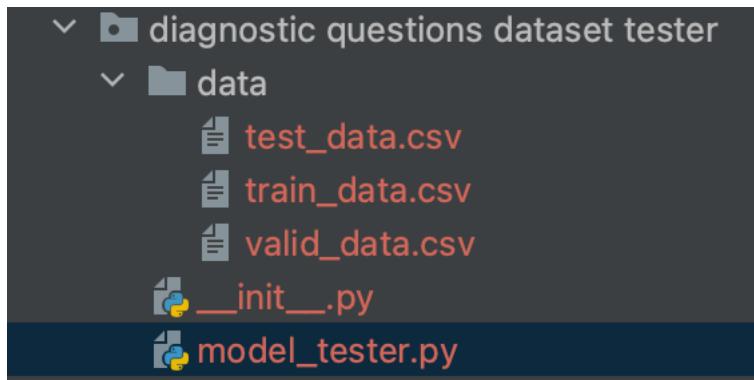
And then trained NMF on the following hyperparameters:

```
nmf = NMF(
    n_components=2,
    alpha_W=20,
    alpha_H=0,
    l1_ratio=0.0,
    init='random',
    solver='mu',
    beta_loss='kullback-leibler',
    max_iter=10000,
    tol=0.001,
    random_state=22
)
```

In the case of missing values the next hyperparameters have to take the values:

- solver: 'mu'
- init method: 'random'

In PyCharm, I created a module with the next structure:



Where `model_tester.py` has the code shown below

```
import sklearn
from sklearn.decomposition import NMF

from sklearn.metrics import accuracy_score
from scipy.sparse import csc_array

import numpy as np
import pandas as pd

print(sklearn.__version__)

# _____Retrieve data in required format_____

def get_data_path(data_type: str):
    return 'data/' + data_type + '_data.csv'

def download_data(data_type: str):
    data_path = get_data_path(data_type)
    data = pd.read_csv(data_path)
    return data

def convert_to_dataframe(data):
    df = pd.DataFrame(
        zip(data['user_id'], data['question_id'], data['is_correct'])
    )
    df.columns = ['user_id', 'question_id', 'is_correct']
    return df

def convert_to_array_keeping_nan(data):
    mask = csc_array(([1] * len(data['is_correct']), (data['user_id'], data['question_id'])))

    mask = mask.toarray().astype(np.float32)
    # now we can mark zero as nan
    mask[mask == 0] = np.nan

    matrix_sparse = csc_array(
        (data.is_correct, (data.user_id, data.question_id))
    )
    data_nan_array = matrix_sparse.toarray() * mask
    return data_nan_array
```

```

data_types = ['train', 'valid', 'test']
given_data = {}

for d_type in data_types:
    given_data[d_type] = convert_to_dataframe(download_data(d_type))

train_valid_df = pd.concat([given_data['train'], given_data['valid']])
train_valid_df.reset_index(drop=True, inplace=True)
given_data['train_valid'] = train_valid_df

train_valid_nan_array = convert_to_array_keeping_nan(given_data['train_valid'])
train_nan_array = convert_to_array_keeping_nan(given_data['train'])

print("train_valid matrix[: 10, : 10]")
print(train_valid_nan_array[: 10, : 10])
print(train_valid_nan_array.shape)

# _____Predict student-question interaction matrix_____

```

---

```

nmf = NMF(
    n_components=2,
    alpha_W=20,
    alpha_H=0,
    l1_ratio=0.0,
    init='random',
    solver='mu',
    beta_loss='kullback-leibler',
    max_iter=100000,
    tol=0.001,
    random_state=22
)

```

---

```

def convert_to_binary(matrix):
    matrix = np.where(matrix - 0.5 > 0, 1, 0)
    return matrix

```

---

```

def get_predicted_matrix(data):
    W = nmf.fit_transform(data)
    H = nmf.components_

    reconstructed_matrix = np.dot(W, H)
    predicted_matrix = convert_to_binary(reconstructed_matrix)
    return predicted_matrix

```

---

```

predicted_matrixes = {
    'train': get_predicted_matrix(train_nan_array),
    'train_valid': get_predicted_matrix(train_valid_nan_array)
}

```

---

```

# _____Test predictions_____

```

---

```

real_values = {}
for d_type, df in given_data.items():
    real_values[d_type] = df['is_correct']

real_values['train_valid'] = pd.concat([real_values['train'], real_values['valid']], ignore_index = True)

```

---

```

def get_predictions(data, matrix):
    predictions = []

    for u, q in zip(data['user_id'], data['question_id']):
        predictions.append(matrix[u, q])

```

```
    return predictions

def get_score(real_data, predictions):
    accuracy = accuracy_score(real_data, predictions)
    return accuracy

# _____ Scores of predictions from train dataset _____
print('Scores of predictions from nmf trained on train dataset')

train_scores_nmf = get_score(
    real_values['train'],
    get_predictions(given_data['train'], predicted_matrixes['train'])
)
valid_scores_nmf = get_score(
    real_values['valid'],
    get_predictions(given_data['valid'], predicted_matrixes['train'])
)

test_scores_nmf = get_score(
    real_values['test'],
    get_predictions(given_data['test'], predicted_matrixes['train'])
)

scores = []
column_names = []

scores.append(train_scores_nmf)
scores.append(valid_scores_nmf)
scores.append(test_scores_nmf)
column_names.extend(['train', 'valid', 'test'])

scores_df = pd.DataFrame.from_dict(
    scores
)
scores_df = scores_df.T
scores_df.columns = column_names

print(scores_df)

# _____ Scores of predictions from train_valid dataset _____
print('Scores of predictions from nmf trained on train_valid dataset')

train_valid_scores_nmf = get_score(
    real_values['train_valid'],
    get_predictions(given_data['train_valid'], predicted_matrixes['train_valid'])
)

test_scores_nmf = get_score(
    real_values['test'],
    get_predictions(given_data['test'], predicted_matrixes['train_valid'])
)

scores = []
column_names = []

scores.append(train_valid_scores_nmf)
scores.append(test_scores_nmf)
column_names.extend(['train_valid', 'test'])

scores_df = pd.DataFrame.from_dict(
    scores
)
scores_df = scores_df.T
scores_df.columns = column_names
```

```
print(scores_df)
```

And after training the model I received such results:

```
train_valid matrix[: 10, : 10]
[[ 0. nan nan nan nan nan nan nan nan nan]
 [nan  0. nan nan nan nan nan nan nan nan]
 [nan nan  1. nan nan nan nan nan nan nan]
 [nan nan  0.  1. nan nan nan nan nan nan]
 [nan nan nan nan  0. nan nan nan nan nan]
 [nan  1. nan nan nan  0. nan  0. nan nan]
 [nan nan nan nan  0. nan  1. nan nan nan]
 [nan nan nan nan nan nan nan  1. nan  1.]
 [nan nan nan nan nan nan nan nan nan nan]
 [ 0. nan nan nan nan nan nan nan nan  1.]]
(542, 1774)

Scores of predictions from nmf trained on train dataset
      train      valid      test
0  0.751835  0.693903  0.695456

Scores of predictions from nmf trained on train_valid dataset
      train_valid      test
0       0.748957  0.703076
```

From this results we see that the testing score improves when we use more data for training. It makes sense because the algorithm focuses only on the present non nan values and we have only 5.9% of the entries filled with values

This model trained on the combined train and valid datasets gives the best performance over all the models

## 8. Summary

The best result all models produce with number of factors 2

NMF model achieves the highest score comparing to the models of other algorithms:

- 0.695456 accuracy on the algorithm that doesn't accept missing values
- 0.703076 accuracy on the algorithm that accepts missing values

For the NMF algorithm that accepts missing values the size of data and the amount of non nan values in data play an important role

On all algorithms, hyperparameters Optimization with cross validation converges to some consistent result after a certain number of iterations

## 9. Future steps

1. To implement ranking step using Matrix Factorization algorithms with ranking metric like MAP@K for recommending next diagnostic question to reveal students' misconceptions

2. To group data first and then apply Matrix Factorization Algorithms to each group. We can group data by types of questions taken from questions metadata, or by the result from clustering. For instance, this [notebook](#) performs clustering of the data. If the number of factors represent the number of groups of students and questions then with factor value 2 we probably capture only the high level features but not the variety of concepts in the learning material. With splitting by groups we could try to reveal more interesting subgroups that cover concepts from learning material
3. To try ensemble methods, hybrid approaches like these [types of hybrid recommender systems](#)
4. To use Matrix Factorization Algorithms for binary data like [Non-negative Matrix Factorization for Binary Data](#) or [Binary Matrix](#)